

Homework 3

STAT 471

*Jacob Kahn
Devesh Dayal
Meghana Jayam
30 October, 2016*

Preflight Tasks

We have a lot of existing variables, so we'll clean them out.

```
rm(list=ls())
```

Check our working directory. This changed for various group members, so we each set it locally:

```
# setwd(dir)  
# getwd() # check that working directory
```

Problem 1

Part a

Generate a predictor X of length $n=100$, as well as a noise vector ϵ of length $n=100$.

```
set.seed(10)  
X <- rnorm(100)  
epsilon <- rnorm(100)
```

Part b

Generate a response vector Y of length $n=100$ with B_0, B_1, B_2, B_3

```
B0 <- -2  
B1 <- 0.1  
B2 <- 1  
B3 <- 5  
Y <- B0+B1*X+B2*X^2+B3*X^3+epsilon
```

Part c

Perform best subset selection (find C_p , BIC, adjr^2)

```

library(leaps) # for regsubsets
data <- data.frame(y=Y, x=X)
regsub <- regsubsets(y~poly(x,degree=10,raw=TRUE), data=data, nvmax=10)
reg.sum <- summary(regsub)
names(reg.sum)

## [1] "which" "rsq" "rss" "adjr2" "cp" "bic" "outmat" "obj"

# find optimal size by getting min Cp, BIC, adjr2
which.min(reg.sum$cp) # 2

## [1] 2

which.min(reg.sum$bic) # 2


## [1] 2

which.max(reg.sum$adjr2) # 2


## [1] 2

# plot Cp, BIC, adjr2
plot(reg.sum$cp, xlab="Number of predictors", ylab="Cp", col="red", type="p", pch=16) # exhibit 1


```



```
plot(reg.sum$bic, xlab="Number of predictors", ylab="BIC", col="red", type="p", pch=16) # exhibit 2
```



```
plot(reg.sum$adjr2, xlab="Number of predictors", ylab="adjr2", col="red", type="p", pch=16) # exhibit 3
```



As seen by the plots, we would use a 2-variable model with Cp, a 2-variable model with BIC, and a 2-variable model with Adjusted R^2 .

```
coef(regsub, id=2)
```

```
##               (Intercept) poly(x, degree = 10, raw = TRUE)2
##               -2.0710776                                0.9642645
## poly(x, degree = 10, raw = TRUE)3
##               5.0163042
```

```
## (Intercept) poly(x, degree = 10, raw = TRUE)2 poly(x, degree = 10, raw = TRUE)3
## -2.0710776                                0.9642645                                5.0163042
```

The model will use x^2 and x^3 .

Part d

Use forward selection and the backward selection to compare results with part c. First, we will do forward selection:

```
regsub.f <- regsubsets(y~poly(x,degree=10,raw=TRUE), data=data, nvmax=10, method="forward")
reg.sum.f <- summary(regsub.f)

which.min(reg.sum.f$cp)      # 2
```

```
## [1] 2
```


```
which.min(reg.sum.f$bic)      # 2
```

```
## [1] 2
```


```
which.max(reg.sum.f$adjr2)    # 2
```

```
## [1] 2
```


```
plot(reg.sum.f$cp, xlab="Number of predictors", ylab="Cp", col="red", type="p", pch=16) # exhibit 4
```



```
plot(reg.sum.f$bic, xlab="Number of predictors", ylab="BIC", col="red", type="p", pch=16) # exhibit 5
```



```
plot(reg.sum.f$adjr2, xlab="Number of predictors", ylab="adjr2", col="red", type="p", pch=16) # exhibit
```



```
coef(regsub.f, id=2)
```

```
##              (Intercept) poly(x, degree = 10, raw = TRUE)2
##              -2.0710776                                0.9642645
## poly(x, degree = 10, raw = TRUE)3
##              5.0163042
```

```
## (Intercept) poly(x, degree = 10, raw = TRUE)2 poly(x, degree = 10, raw = TRUE)3
## -2.0710776                                0.9642645                                5.0163042
```

For forward selection, we see the same results as in part c. Now, we will do backward selection:

```
regsub.b <- regsubsets(y~poly(x,degree=10,raw=TRUE), data=data, nvmax=10, method="backward")
reg.sum.b <- summary(regsub.b)
```

```
which.min(reg.sum.b$cp)      # 2
```

```
## [1] 2
```


```
which.min(reg.sum.b$bic)     # 2
```

```
## [1] 2
```


```
which.max(reg.sum.b$adjr2)    # 2
```

```
## [1] 2
```


```
plot(reg.sum.b$cp, xlab="Number of predictors", ylab="Cp", col="red", type="p", pch=16) # exhibit 7
```



```
plot(reg.sum.f$bic, xlab="Number of predictors", ylab="BIC", col="red", type="p", pch=16) # exhibit 8
```

```
plot(reg.sum.f$adjr2, xlab="Number of predictors", ylab="adjr2", col="red", type="p", pch=16) # exhibit
```



```
coef(regsub.b, id=2)
```

```
##                (Intercept) poly(x, degree = 10, raw = TRUE)2
##                -2.0710776                                0.9642645
## poly(x, degree = 10, raw = TRUE)3
##                5.0163042
```

```
## (Intercept) poly(x, degree = 10, raw = TRUE)2 poly(x, degree = 10, raw = TRUE)3
## -2.0710776                                0.9642645                                5.0163042
```

Again, we see the same model and values.


Part e

Fit a LASSO model and use cross-validation to select the optimal value for lambda.


```
library(glmnet) # for LASSO
```

```
## Loading required package: Matrix
## Loading required package: foreach
## Loaded glmnet 2.0-5
```

```
X.lasso <- model.matrix(y~poly(x,degree=10,raw=TRUE), data=data)[-1]
Y.lasso <- Y
lasso.cv <- cv.glmnet(X.lasso, Y.lasso, alpha=1, nfolds=10)
plot(lasso.cv) # exhibit 10
```



```
plot(lasso.cv$lambda) # exhibit 11
```



```
lambda.final <- lasso.cv$lambda.min # 0.05631109
```

```
beta <- coef(lasso.cv, s=lambda.final)
```

```
beta <- as.matrix(beta)
```

```
beta
```

```
##                                     1
## (Intercept)                       -1.9625901145
## poly(x, degree = 10, raw = TRUE)1  0.0056439028
## poly(x, degree = 10, raw = TRUE)2  0.7397802716
## poly(x, degree = 10, raw = TRUE)3  4.9391488225
## poly(x, degree = 10, raw = TRUE)4  0.0361210858
## poly(x, degree = 10, raw = TRUE)5  0.0041259611
## poly(x, degree = 10, raw = TRUE)6  0.0000000000
## poly(x, degree = 10, raw = TRUE)7  0.0000000000
## poly(x, degree = 10, raw = TRUE)8  0.0000000000
## poly(x, degree = 10, raw = TRUE)9  0.0003331596
## poly(x, degree = 10, raw = TRUE)10 0.0000000000
```

```
#                                     1
# (Intercept)                       -1.9625901145
# poly(x, degree = 10, raw = TRUE)1  0.0056439028
# poly(x, degree = 10, raw = TRUE)2  0.7397802716
# poly(x, degree = 10, raw = TRUE)3  4.9391488225
# poly(x, degree = 10, raw = TRUE)4  0.0361210858
```

```
# poly(x, degree = 10, raw = TRUE)5 0.0041259611
# poly(x, degree = 10, raw = TRUE)6 0.0000000000
# poly(x, degree = 10, raw = TRUE)7 0.0000000000
# poly(x, degree = 10, raw = TRUE)8 0.0000000000
# poly(x, degree = 10, raw = TRUE)9 0.0003331596
# poly(x, degree = 10, raw = TRUE)10 0.0000000000
```

Using the LASSO method, the model picks x , x^2 , x^3 , x^4 , x^5 , x^9 . The coefficients for x^9 , x , and x^5 are more negligible than the others.

Part f

```
B7 <- 7
Y.f <- B0 + B7*X^7 + epsilon
data.f <- data.frame(y=Y.f, x=X)
```

Model selection using regsubsets:

```
regsub.partf <- regsubsets(y~poly(x,degree=10,raw=TRUE), data=data.f, nvmax=10)
reg.sum.partf <- summary(regsub.partf)
```

```
which.min(reg.sum.partf$cp)      # 1
```

```
## [1] 1
```

```
which.min(reg.sum.partf$bic)     # 1
```

```
## [1] 1
```

```
which.max(reg.sum.partf$adjr2)   # 1
```

```
## [1] 1
```


```
coef(regsub.partf, id=1)
```

```
##                (Intercept) poly(x, degree = 10, raw = TRUE)7
##                -2.095239                      6.999867
```

```
# (Intercept) poly(x, degree = 10, raw = TRUE)7
# -2.095239                      6.999867
```

All three methods chose one-variable models. Now, using LASSO:

```
X.lasso.f <- model.matrix(y~poly(x,degree=10,raw=TRUE), data=data.f)[,-1]
Y.lasso.f <- Y.f
lasso.cv.f <- cv.glmnet(X.lasso.f, Y.lasso.f, alpha=1, nfolds=10)
plot(lasso.cv.f) # exhibit 12
```



```
lambda.final.f <- lasso.cv.f$lambda.min # 10.43878
```

```
beta <- coef(lasso.cv.f, s=lambda.final.f)
```

```
beta <- as.matrix(beta)
```

```
beta
```

```
##                                1
## (Intercept)                   -2.561234
## poly(x, degree = 10, raw = TRUE)1  0.000000
## poly(x, degree = 10, raw = TRUE)2  0.000000
## poly(x, degree = 10, raw = TRUE)3  0.000000
## poly(x, degree = 10, raw = TRUE)4  0.000000
## poly(x, degree = 10, raw = TRUE)5  0.000000
## poly(x, degree = 10, raw = TRUE)6  0.000000
## poly(x, degree = 10, raw = TRUE)7  6.775923
## poly(x, degree = 10, raw = TRUE)8  0.000000
## poly(x, degree = 10, raw = TRUE)9  0.000000
## poly(x, degree = 10, raw = TRUE)10 0.000000
```

```
#                                1
# (Intercept)                   -2.561234
# poly(x, degree = 10, raw = TRUE)1  0.000000
# poly(x, degree = 10, raw = TRUE)2  0.000000
# poly(x, degree = 10, raw = TRUE)3  0.000000
# poly(x, degree = 10, raw = TRUE)4  0.000000
```

```
# poly(x, degree = 10, raw = TRUE)5 0.000000
# poly(x, degree = 10, raw = TRUE)6 0.000000
# poly(x, degree = 10, raw = TRUE)7 6.775923
# poly(x, degree = 10, raw = TRUE)8 0.000000
# poly(x, degree = 10, raw = TRUE)9 0.000000
# poly(x, degree = 10, raw = TRUE)10 0.000000
```

Using LASSO, the model selected is also a one-variable model.

Problem 2